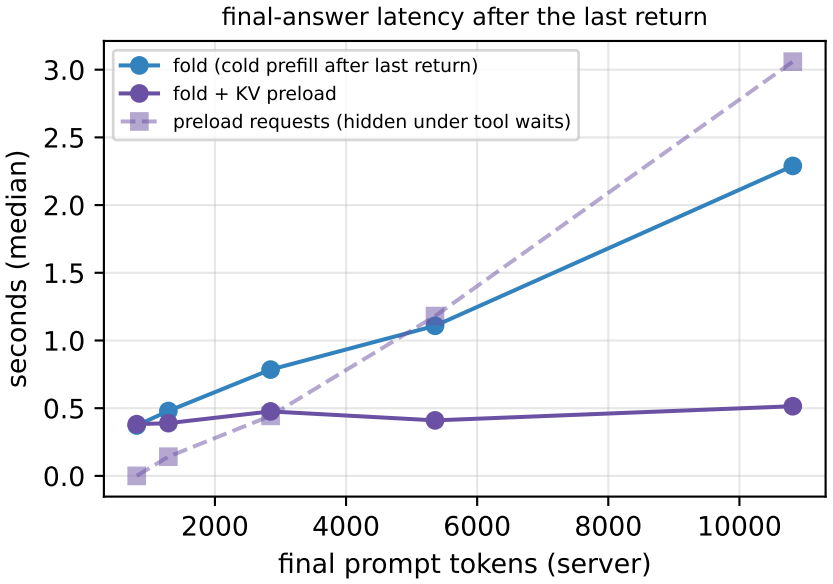
final-answer latency after the last return
fold (cold prefill after last return)
fold + KV preload
preload requests (hidden under tool waits)
seconds (median)
0.0
0.5
1.0
1.5
2.0
2.5
3.0
2000
4000
6000
8000
10000
final prompt tokens (server)

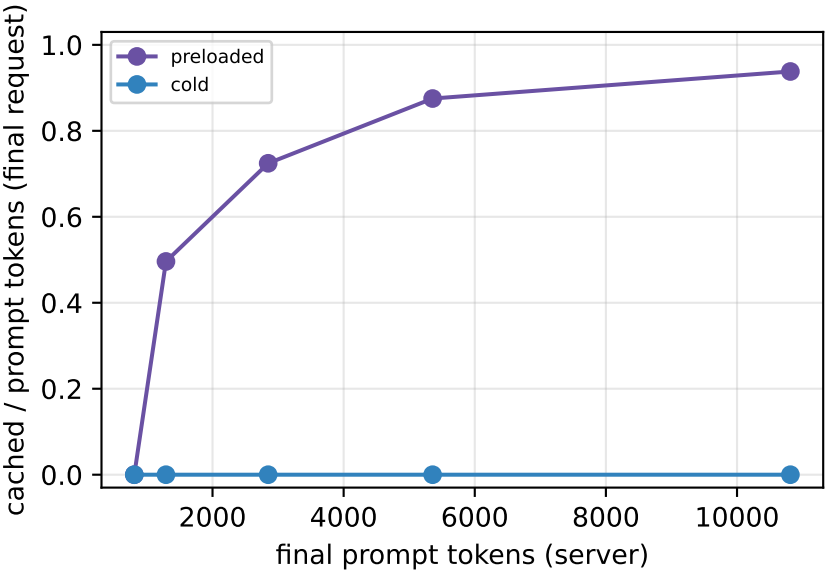
preloaded
cold
cached / prompt tokens (final request)
0.0
0.2
0.4
0.6
0.8
1.0
2000
4000
6000
8000
10000
final prompt tokens (server)