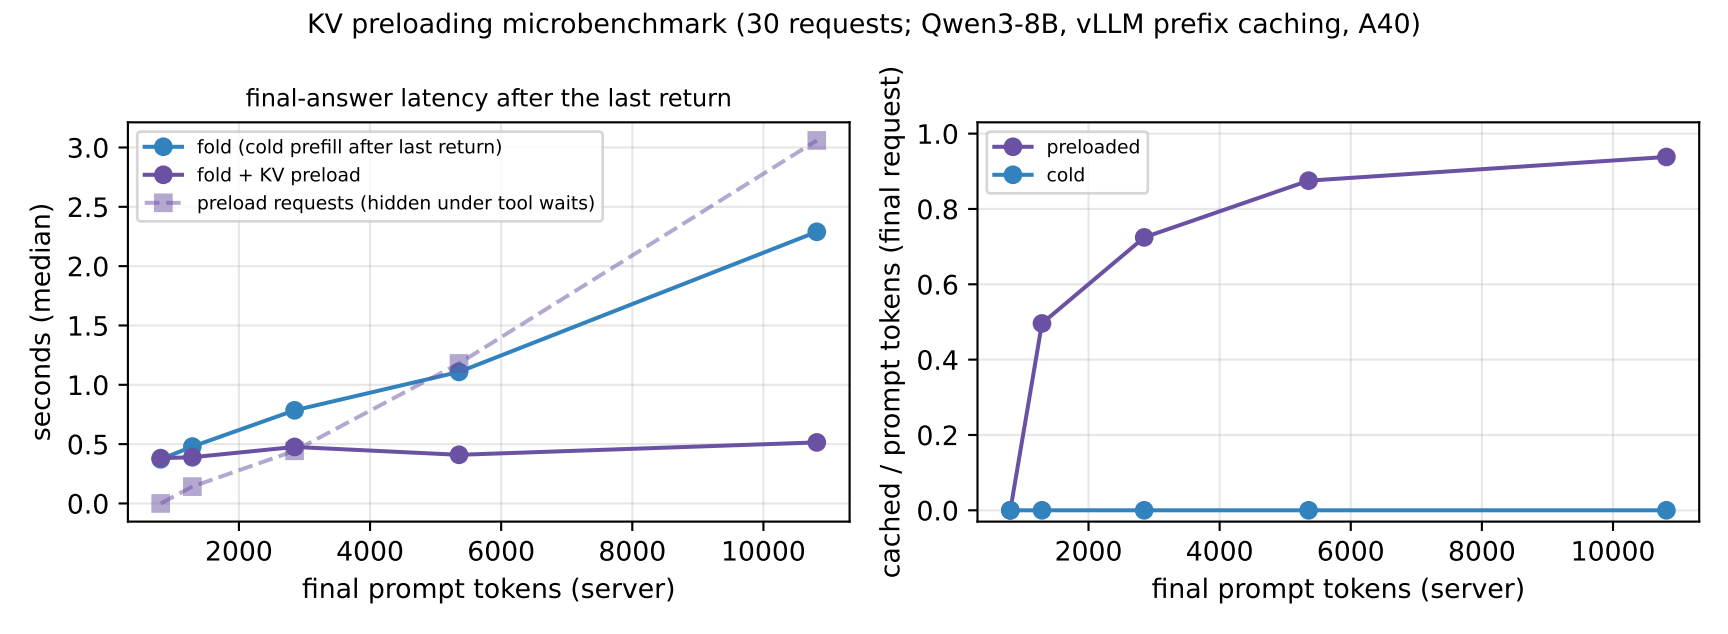
KV preloading microbenchmark (30 requests; Qwen3-8B, vLLM prefix caching, A40)
final-answer latency after the last return
fold (cold prefill after last return)
fold + KV preload
preload requests (hidden under tool waits)
seconds (median)
final prompt tokens (server)
preloaded
cold
cached / prompt tokens (final request)
final prompt tokens (server)